

PRODUCT & VISION • 2026

kodwai

*Where developers prove how they work with AI
agents, and companies hire for it.*

Live

IN PRODUCTION

Free

FOR DEVELOPERS

Built

CLI • WEB • API

THE WEDGE • EVERY COMPANY HIRES FOR AI-AGENT SKILL, AND NO PLATFORM MEASURES IT

81%

of Big Tech interviewers suspect AI cheating

31% have caught a candidate. The legacy LeetCode interview no longer signals skill. (interviewing.io, 2025)

84%

of developers use or plan to use AI tools

51% use them every day. Directing an agent has become the job. (Stack Overflow survey, 2025)

29%

of developers trust AI output accuracy

The scarce, hireable skill is the human judgment to verify and steer agents. (Stack Overflow, 2025)

TWO PLATFORMS. ONE SKILL.

Developer-first. *Business-funded.*

DEVELOPER PLATFORM · FREE

Solve real, ticket-sized challenges with Claude Code, Cursor, or Codex, on your own machine. Get scored on how well you direct it. Climb a public leaderboard.

- Free for every developer, forever
- The audience and the data moat
- The top of the funnel

BUSINESS PLATFORM · PAID

Interview on deeper, role-specific challenges, or author your own. Watch the real process, not just the answer, then score and compare candidates.

- ◆ Sold to hiring teams and companies
- ◆ Where the revenue comes from
- ◆ Companies bring their own API key, so margins run ~85–90%

FLYWHEEL Free developers build the audience and the labeled dataset. We then sell the hiring product to the companies those developers already work at.

THE AI COLLABORATION SCORE · FREE, LIVE, IN PRODUCTION

Not a pass/fail. *A complete picture.*

Direction · 6 signals · 50 pts

88%

Outcome · 3 signals · 35 pts

92%

Lift · 2 signals · 15 pts

74%

OVERALL SCORE

86

/100

Solve via `npx @kodwai/cli challenge <slug>`, then 11 signals across three axes: algorithmic where it can be (tests, code quality, complexity), LLM-judged where it must be (spec precision, verification, decomposition, recovery). Per-signal evidence, exportable.

Hiring, measured.

You stop grading the *answer*. You start watching the work.

Invite a candidate by link. They solve on their own machine with a real agent, and you watch the real process live, not just the final answer.

- 01 Solve on their own machine with a real agent
- 02 Watch every prompt, commit, and test, live
- 03 Observer dashboard: diffs, tool feed, timer, cost
- 04 Dual scored, AI + human · AES-256, HMAC events

session.log · candidate_4821

transcript ▶ "refactor the auth guard, keep the suite green"

```
41  const score = run(tests)
42  - return score
42  + return clamp(score, 0, 100)
```

00:00.4 ● clone repo 412ms

00:01.1 ● run test suite 4/4 ✓ 1.4s

00:02.6 ● analyze coverage...

ELAPSED 04:18 TOKENS 12.4k COST \$0.038

AI 87 · HUMAN +A
dual scorecard

87
STRONG HIRE

BUSINESS MODEL

Free for developers. *Recurring* for companies.

The candidate brings the API key, so the inference bill is theirs, not ours. That zero cost of goods is the whole margin story.

\$0 LLM COST · ~85-90% MARGIN · MODELED

~/kodwai · pricing

\$ kodwai pricing --tiers --econ

DEVELOPER

developer	free	// forever
developer.pro	future	// company-tagged banks

BUSINESS

starter	\$60/mo	5 seats · 50 sessions
▶ growth	\$200/mo	✓ recommended
enterprise	custom	→ \$15-25K ACV

> computing unit economics ...

cogs_llm	\$0.00	◀ customers bring their own key
----------	--------	---------------------------------

GROSS MARGIN

~85-90%

✓ modeled

\$0 inference cost to Kodwai

MARKET OPPORTUNITY

A massive market with an *empty quadrant*.



TAM ~\$1.4B Coding and technical assessment software, 2024 (DataIntello, Verified Market Reports)

SAM \$280-420M AI-native assessment: the 20-30% of teams measuring agent skill (Kodwai estimate)

SOM \$3-8M Year 2-3 ARR: 150-400 paying teams at \$15-25K ACV (target)

GROWTH ~12-15% CAGR toward ~\$3-4.5B by 2030. The agent-native, own-terminal, free-community quadrant is still empty (mid-2026).

The signals are *loud*.

Oct 2025

Meta ships an AI-enabled coding interview

Framed as more representative of the environment its future employees will work in. The biggest employer concedes the premise. (Hello Interview)

2025-26

Autonomous coding agents went mainstream

Claude Code at a ~\$2.5B run rate, authoring ~4% of public GitHub commits; Cursor past \$2B ARR. Directing agents is the job. (SemiAnalysis)

Now

Incumbents validated the category, stayed in the browser

CodeSignal and HackerRank shipped agentic assessments, but none paired a real agent in the developer's own terminal with a free, terminal-native community. That intersection is empty.

Moat

The moat: a dataset no incumbent can buy

The Agent SDK is public. The edge is the community and the real, labeled agent-collaboration data it produces. Category ownership, won in a 12-18 month window.

TEAM

Built by people who *live this problem.*



Ege Hakan Karaagac

Co-Founder & CEO

5 years building scalable backend systems and AI products. Leads product across the full stack: CLI, web app, API, and the AI scoring pipeline.



Dogukan Ertunga Kurnaz

Co-Founder & CTO

4 years in DevSecOps, infrastructure security, and automation. Leads Kodwai's cloud infrastructure, CI/CD, and security hardening across all services.



Two people shipped the CLI, web app, API, and scoring engine, live in production, before taking a dollar. Where developers prove it, companies hire it.
· hakan@kodwai.com · kodwai.com